

Deep-Learning-Series-2-Important-Topics

? For more notes visit

https://rtpnotes.vercel.app

- Deep-Learning-Series-2-Important-Topics
 - 1. Recurrent Neural Network
 - RNN Diagram
 - Computational Graph
 - Applications of RNN
 - Problems in RNN
 - 1. Vanishing Gradient Problem
 - 2. Exploding Gradient Problem
 - 3. Short-Term Memory
 - How to Fix These Problems?
 - 1. Use LSTM and GRU
 - 2. Gradient Clipping
 - 3. Residual Connections
 - 4. Use Attention Mechanisms
 - 2. Explain LSTM and GRU in Detail
 - Long Short-Term Memory (LSTM)
 - 1. Forget Gate Decides what information to erase
 - 2. Input Gate Decides what new information to store
 - 3. Output Gate Decides what information to send as output
 - 2) Gated Recurrent Unit (GRU)
 - 1. Update Gate (z) Decides how much old information to keep
 - 2. Reset Gate (r) Decides how much past information to ignore
 - 3. Encoder-Decoder Architecture with an Example
 - Why is Encoder-Decoder Needed?



- How Does Encoder-Decoder Work?
- Step-by-Step Working
- Example: Machine Translation (English to French)
 - 1. Encoding Phase (Encoder RNN)
 - 2. Decoding Phase (Decoder RNN)
- 4. Explain Recursive Neural Network
 - How Does a Recursive Neural Network Work?
 - Structure of Recursive Neural Network
 - Applications of Recursive Neural Networks
 - 1. Natural Language Processing (NLP)
 - 2. Computer Vision
 - 3. Knowledge Graphs
 - Difference Between RNN and RvNN
 - Limitations of Recursive Neural Networks
- 5. Explain CNN
 - Why CNN?
 - CNN Architecture
 - Step-by-Step CNN Processing
 - 3) Example of CNN Architecture
 - Example: Cat vs. Dog Image Classification
- 6. Explain Deep Learning Application
 - Applications of Deep Learning
 - 1. Image Processing and Computer Vision
 - 2. Natural Language Processing (NLP)
 - 3. Healthcare and Drug Discovery
 - 4. Finance and Stock Market
 - 5. Gaming and Robotics
 - Deep Learning Architecture for Machine Translation
 - 1) Traditional vs. Deep Learning-Based Translation
 - 2) Architecture of Deep Learning-Based Machine Translation
 - 3) Step-by-Step Working of Machine Translation Model
- 7. Explain Different Word Embeddings in NLP
 - What is Word Embedding?



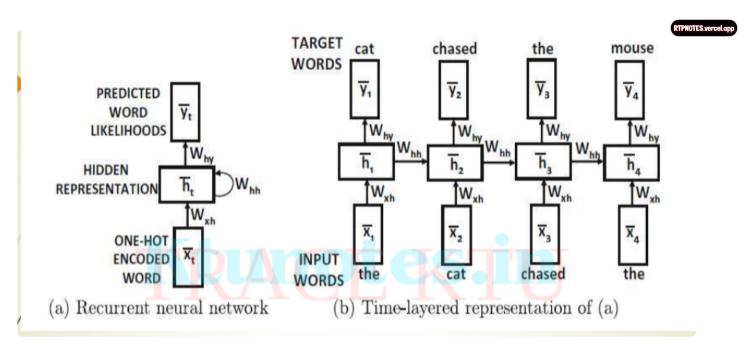
- Why Do We Need Word Embeddings?
- Types of Word Embeddings
- A) Basic Word Embeddings (Simple but Limited)
 - 1. Bag of Words (BoW)
 - 2. Term Frequency-Inverse Document Frequency (TF-IDF)
- B) Advanced Word Embeddings (Smarter & Context-Aware)
 - 3. Word2Vec (Google's Method)
 - 4. GloVe (Stanford's Method)
 - 5. FastText (Facebook's Method)
 - 6. BERT (Google's Transformer Model Used in ChatGPT & Google Search)
- 8. Deep Learning for Speech Processing
 - How Deep Learning Helps in Speech Processing
 - 1. Convolutional Neural Networks (CNNs) → Used for Speech Recognition
 - 2. Recurrent Neural Networks (RNNs) → Used for Speech Understanding
 - 3. Long Short-Term Memory (LSTM) → Used for Better Speech Context
 - 4. Transformers (BERT, Whisper, Wav2Vec) → Used for Real-Time Speech Translation
- 9. Explain Photo Encoders and Transfer Learning
 - How It Works:
 - How Are Photo Encoders Used?
 - What is Transfer Learning?
 - Why is Transfer Learning Useful?
 - How Transfer Learning Works in Images?
 - Where Is Transfer Learning Used?

1. Recurrent Neural Network

RNNs are neural networks designed for **sequential data** like text, time-series, and speech. Unlike regular neural networks where each input is processed independently, RNNs maintain a **hidden state**, allowing them to remember past information while processing new inputs.

Example: While reading a sentence, we understand each word based on the previous words.

RNN Diagram



An RNN processes inputs step by step while maintaining a hidden state:

$$x1 \rightarrow [h1] \rightarrow x2 \rightarrow [h2] \rightarrow x3 \rightarrow [h3] \rightarrow y3$$

- x = Input at each time step
- h = Hidden state (memory)
- **y** = Output

The hidden state helps retain past information, making RNNs useful for tasks like **language modeling and speech recognition**.

The hidden state at time t is given by a function of the input vector at time t and the hidden vector at time (t-1):

$$\overline{h}_t = f(\overline{h}_{t-1}, \overline{x}_t)$$

Computational Graph

The computational graph of an RNN represents how the network processes inputs over time.

- When an RNN is **unrolled**, it looks like a deep neural network.
- Weights are shared across time steps.
- Training uses **Backpropagation Through Time (BPTT)** to update weights.

RTPNOTES.vercel.app

Example: Processing a sentence word by word while maintaining context from previous words.

Applications of RNN

- Google Translate Translates text from one language to another.
- Speech Recognition Used in virtual assistants like Siri and Google Assistant.
- Chatbots Helps AI understand human conversations.
- Stock Market Predictions Analyzes past trends to forecast future prices.

Problems in RNN

1. Vanishing Gradient Problem

 As the network processes long sequences, gradients become too small, making it hard to learn long-term dependencies.

2. Exploding Gradient Problem

• In some cases, gradients become **too large**, making training unstable.

3. Short-Term Memory

RNNs struggle with long sequences, forgetting older inputs quickly.

Example: Trying to recall a book's first chapter after reading 20 chapters.

How to Fix These Problems?

1. Use LSTM and GRU

- LSTM (Long Short-Term Memory) and GRU (Gated Recurrent Units) introduce **gates** that control what to remember and forget.
- These models prevent vanishing gradients and improve memory retention.

2. Gradient Clipping

Limits the gradient size to prevent exploding gradients.

3. Residual Connections

Helps stabilize training by ensuring smooth information flow.

4. Use Attention Mechanisms



- Instead of processing every input equally, the model focuses on the most important parts
 of the sequence.
- Used in machine translation, speech recognition, and modern AI models.



2. Explain LSTM and GRU in Detail

RNNs struggle with **long-term dependencies** due to **vanishing gradients**. To solve this, researchers introduced **LSTM** and **GRU**, which use **gates** to control the flow of information.

Long Short-Term Memory (LSTM)

LSTM is an improved version of **Recurrent Neural Networks (RNNs)** that helps remember important information **for a long time**. It was designed to **solve the problems of vanishing gradients and short-term memory** in simple RNNs.

Simple RNNs struggle with **long sequences** because:

- 1. They forget older information quickly.
- 2. They find it hard to learn from long-term dependencies (vanishing gradient problem).

LSTM fixes these problems using **memory cells and gates** that **control what information is kept, updated, or forgotten**.

LSTMs have a **special memory cell** that decides what to **remember and forget** using three gates:

1. Forget Gate - Decides what information to erase

- If the forget gate value is **0**, information is **deleted**.
- If it is **1**, the information is **kept**.

Example: When cleaning your phone, you decide which photos to **delete (forget)** and which to **keep (remember)**.

2. Input Gate – Decides what new information to store

Controls how much new data should be added to the memory.



• If the input gate is **high**, new information is stored.

Example: When studying, you **choose which new facts to remember** while forgetting unnecessary details.

3. Output Gate – Decides what information to send as output

Controls what information is passed to the next step.

Example: When summarizing a book, you **decide which key points to share** and which details to leave out.

2) Gated Recurrent Unit (GRU)

A GRU (Gated Recurrent Unit) is a type of Recurrent Neural Network (RNN) that helps solve the problems of vanishing gradients and short-term memory seen in simple RNNs. It is similar to LSTM (Long Short-Term Memory) but is simpler and faster to train.

Simple RNNs struggle to remember **long sequences** because:

- 1. They forget older information quickly (short-term memory issue).
- 2. They are hard to train for long sequences (vanishing gradient problem).

GRUs fix these problems by using gates to control how much past information is remembered or forgotten.

Instead of just passing information forward like a simple RNN, a GRU chooses what to keep and what to forget using two special gates:

- 1. Update Gate (z) Decides how much old information to keep
 - If z = 1, the GRU keeps all past information.
 - If z = 0, the GRU forgets past information and only uses the new input.

Example: Imagine a student studying for a test. If the update gate is high, they **remember old concepts**. If it is low, they **focus only on new material**.

- 2. Reset Gate (r) Decides how much past information to ignore
 - If r = 0, the GRU forgets past information completely, like starting fresh.
 - If r = 1, the GRU keeps past information while updating with new input.



Example: When writing an essay, if the reset gate is low, the writer **forgets old ideas** and starts fresh. If it is high, they **build on previous ideas**.



3. Encoder-Decoder Architecture with an Example

The **Encoder-Decoder** architecture is a special type of **Recurrent Neural Network (RNN)** designed to handle **sequence-to-sequence (seq2seq) tasks**, where the **input and output sequences are of different lengths**.

Why is Encoder-Decoder Needed?

- In regular RNNs, the input and output have the same length.
- In many tasks like machine translation, chatbot responses, and summarization, the output length is different from the input length.
- Example: Translating "Hello" (1 word) into "Bonjour, comment ça va?" (4 words).

How Does Encoder-Decoder Work?

The Encoder-Decoder model consists of two RNNs:

- 1. **Encoder** Reads the input sequence and encodes it into a **context vector**.
- 2. **Decoder** Takes the context vector and generates the output sequence.

Step-by-Step Working

Encoder (Input Processing)

- Reads the input word by word (or token by token).
- Produces a context vector (fixed-size representation of the input).

Context Vector

- Stores the summarized information of the entire input sequence.
- Acts as a bridge between the encoder and decoder.

Decoder (Output Generation)

Takes the context vector as input.



Generates the output one word at a time, based on previously generated words.

Example: Machine Translation (English to French)

Let's say we want to translate:

```
Input (English): "I love pizza"
Expected Output (French): "J'adore la pizza"
```

- 1. Encoding Phase (Encoder RNN)
 - The **encoder reads** the English words one by one.
 - It **creates a final hidden state** (context vector) that represents the entire sentence.

```
Input: "I love pizza"
Step 1: h1 = RNN(x1) → "I"
Step 2: h2 = RNN(h1, x2) → "love"
Step 3: h3 = RNN(h2, x3) → "pizza"
Final Context Vector (h3): Summarizes "I love pizza"
```

2. Decoding Phase (Decoder RNN)

 The decoder takes the context vector and generates the French translation word by word.

```
Context Vector → Decoder Starts
Step 1: "J'" (Output: y1)
Step 2: "adore" (Uses y1 as input)
Step 3: "la" (Uses y2 as input)
Step 4: "pizza" (Uses y3 as input)
Final Output: "J'adore la pizza"
```



4. Explain Recursive Neural Network

A **Recursive Neural Network (RvNN)** is a type of neural network that processes data with a **hierarchical structure**, such as **trees or graphs**, instead of a simple sequence like RNNs. It is



different from **Recurrent Neural Networks (RNNs)**, which process sequential data (like sentences or time-series).

- Key Idea: Instead of processing data step by step (like RNNs), RvNN builds a tree-like structure where information from smaller parts is combined to form a larger understanding.
- **Example**: Understanding a sentence based on **grammatical structure** rather than just word order.

How Does a Recursive Neural Network Work?

- Tree-structured processing: RvNN processes data in a hierarchical manner, merging smaller units into larger representations.
- Example (Sentence Parsing):

Consider the sentence:

"(The cat) (sits (on the mat))"

Here, RvNN would first merge:

```
    "The" + "cat" → Phrase: "The cat"
    "on" + "the mat" → Phrase: "on the mat"
    "sits" + "on the mat" → Phrase: "sits on the mat"
    "The cat" + "sits on the mat" → Final representation of the sentence
```

Structure of Recursive Neural Network

```
Sentence
/ \
Phrase 1 Phrase 2
/ \ / \
Word1 Word2 Word3 Word4
```

This hierarchical process allows the network to **understand relationships** between words beyond just their order.

Applications of Recursive Neural Networks

1. Natural Language Processing (NLP)



- Sentence Parsing: Understanding grammar structures.
- Sentiment Analysis: Understanding sentiment based on hierarchical phrases instead of just word order.

2. Computer Vision

- Scene Understanding: Understanding images in a part-whole hierarchy.
- Example: Recognizing a **face** from eyes, nose, and mouth.

3. Knowledge Graphs

 Used in question answering systems where knowledge is stored in graphs instead of simple lists.

Difference Between RNN and RvNN

Feature	Recursive Neural Network (RvNN)	Recurrent Neural Network (RNN)
Data Structure	Tree-like structure	Sequential structure
Processing Order	Bottom-up (from small parts to big concepts)	Left-to-right (or time-step based)
Best Used For	NLP, hierarchical data, parsing, graphs	Speech, time-series, machine translation
Example	Understanding sentence structure	Predicting the next word in a sentence

Limitations of Recursive Neural Networks

1. Requires Predefined Structure

- Unlike RNNs, RvNNs need a tree structure before training.
- Example: In NLP, it needs a predefined sentence structure.

2. Computational Complexity

• Since it **builds trees**, it is more computationally expensive than RNNs.

3. Not Suitable for All Sequential Data



If the data does not have a clear hierarchy, RvNNs may not be effective.



5. Explain CNN

A Convolutional Neural Network (CNN) is a deep learning model used primarily for image recognition, object detection, and pattern recognition. Unlike regular neural networks, which process all pixels independently, CNNs capture spatial relationships in images through specialized layers.

Why CNN?

- Traditional neural networks (fully connected layers) treat every pixel as independent, leading to huge numbers of parameters.
- CNNs reduce parameters by using convolution operations to detect features (edges, textures, objects) efficiently.
- Instead of manually designing feature detectors, CNNs learn them automatically through training.

CNN Architecture

A CNN typically consists of **four main layers**:

- 1. Convolution Layer (Feature Extraction)
- 2. Pooling Layer (Downsampling)
- 3. Fully Connected Layer (Classification)
- 4. Activation Functions (ReLU, Softmax)

Step-by-Step CNN Processing

- Step 1: Convolution Layer
- Applies filters (kernels) that slide over the image to extract features (edges, textures, etc.).
- Each filter detects different patterns in the image.
- The output of this layer is called a **Feature Map**.
- Step 2: Activation Function (ReLU)



- ReLU (Rectified Linear Unit) replaces negative values with zero, introducing nonlinearity.
- Helps in detecting complex patterns in images.
- Step 3: Pooling Layer (Downsampling)
- Reduces the size of feature maps while keeping important information.
- Types of pooling:
 - Max Pooling → Takes the largest value in a region.
 - Average Pooling → Takes the average value in a region.
- Step 4: Fully Connected Layer
- Converts the extracted features into a **final classification** (e.g., "Cat" vs. "Dog").
- Uses a softmax activation function to generate class probabilities.

3) Example of CNN Architecture

Let's classify images of cats and dogs using CNN:

Example: Cat vs. Dog Image Classification

- 1. **Input Image:** 32×32 RGB image of a cat.
- Convolution Layer 1: Detects edges and basic patterns.
- 3. **Pooling Layer 1:** Reduces image size while preserving features.
- 4. Convolution Layer 2: Detects more complex features like eyes, whiskers.
- 5. **Pooling Layer 2:** Further reduces image size.
- 6. **Fully Connected Layer:** Converts extracted features into a classification.
- 7. **Softmax Layer:** Outputs probability for each class (Cat = 90%, Dog = 10%).



6. Explain Deep Learning Application

Deep Learning is a subset of **Machine Learning** that uses **neural networks with multiple layers** to automatically learn patterns from data. Unlike traditional models, Deep Learning can **handle large datasets and complex problems** without manual feature engineering.

Applications of Deep Learning



1. Image Processing and Computer Vision

- ✓ Facial Recognition → Used in phones and security systems.
- ✓ Medical Imaging → Detecting diseases from X-rays, MRIs.
- ✓ **Self-driving Cars** → Identifying objects, lane detection.
- 2. Natural Language Processing (NLP)
- ✓ Machine Translation → Google Translate.
- ✓ Chatbots and Virtual Assistants
 → Alexa, Siri, ChatGPT.
- ✓ Speech Recognition → Voice-to-text applications.
- 3. Healthcare and Drug Discovery
- ✓ Disease Prediction → AI diagnosing cancer, diabetes.
- ✓ Protein Structure Prediction → Used in drug discovery.
- 4. Finance and Stock Market
- ✓ Fraud Detection → Identifying fraudulent transactions.
- ✓ Stock Price Prediction → Using historical data for forecasting.
- 5. Gaming and Robotics
- ✓ AI in Video Games → Realistic NPC behaviors.
- ✓ **Industrial Robots** → Used in manufacturing and automation.

Deep Learning Architecture for Machine Translation

Machine Translation refers to the automatic **conversion of text from one language to another** using deep learning.

- 1) Traditional vs. Deep Learning-Based Translation
 - Rule-Based & Statistical Models: Used in early translation systems but lacked fluency.
 - **Deep Learning Approach**: Uses **Neural Machine Translation (NMT)**, which improves translation **accuracy and context understanding**.
- 2) Architecture of Deep Learning-Based Machine Translation

The most common architecture for machine translation is **Sequence-to-Sequence (Seq2Seq) Model with Attention Mechanism**.

RTPNOTES.vercel.app

Encoder-Decoder Model

- Encoder processes the input sentence and creates a context vector.
- Decoder generates the translated output sentence.
- Uses LSTM/GRU for better long-term memory handling.

Attention Mechanism

- Instead of using a fixed-size context vector, attention allows the model to focus on important words in the input sentence during translation.
- Helps in translating long sentences accurately.

3) Step-by-Step Working of Machine Translation Model

• **Example**: Translating "I love deep learning" (English) into "J'adore l'apprentissage profond" (French).

Encoder (Input Processing)

- Reads "I love deep learning" word by word.
- Produces a context vector summarizing the sentence.

Attention Mechanism (Focusing on Important Words)

- Instead of treating all words equally, it assigns higher importance to specific words.
- Example: While translating "deep learning", the model focuses more on "apprentissage profond".

Decoder (Output Generation)

 Takes the context vector and generates "J'adore l'apprentissage profond" word by word.



7. Explain Different Word Embeddings in NLP

What is Word Embedding?

In **Natural Language Processing (NLP)**, word embeddings are techniques to represent words as **numbers (vectors)** so that computers can understand them.



Imagine you have the words "King", "Queen", "Apple", and "Orange". A good word embedding will:

- Place **similar words close to each other** (e.g., "Apple" and "Orange" because both are fruits).
- Keep different words far apart (e.g., "King" and "Apple" because they have no relationship).



Why Do We Need Word Embeddings?

Earlier, we used **one-hot encoding** to represent words as numbers, like this:

Word	Encoding (Vector)	
King	100000	
Queen	010000	
Apple	001000	
Orange	000100	

Problem: This method does **not show relationships** between words. "King" and "Queen" are similar but their encodings have no connection.

Solution: Word embeddings create meaningful numerical representations, where similar words have similar numbers.

Types of Word Embeddings

We use different methods to create word embeddings. They are divided into two categories:

- 1. Basic Methods (Traditional Approaches)
 - Bag of Words (BoW)
 - Term Frequency-Inverse Document Frequency (TF-IDF)
- 2. Advanced Methods (Deep Learning Approaches)
 - Word2Vec
 - GloVe
 - FastText
 - BERT (used in ChatGPT, Google Search)



A) Basic Word Embeddings (Simple but Limited)

1. Bag of Words (BoW)

How it works: Counts how many times each word appears in a document.

Example:

If we have two sentences:

🔟 "I love deep learning."

"I love machine learning."

Word	Sentence 1	Sentence 2
1	1	1
love	1	1
deep	1	0
learning	1	1
machine	0	1

Problem: Ignores meaning and context. "Apple" and "Orange" may be similar but BoW **doesn't know that**.



2. Term Frequency-Inverse Document Frequency (TF-IDF)

How it works: Instead of just counting words, **TF-IDF gives more importance to rare words** (e.g., "quantum physics") and less importance to common words (e.g., "the", "is").

Example:

- "I love deep learning" (common words get low importance).
- "Einstein's theory of relativity" (rare words like "relativity" get high importance).

Problem: Still does not understand relationships between words.

Solution: We use **deep learning-based word embeddings**.

B) Advanced Word Embeddings (Smarter & Context-Aware)

3. Word2Vec (Google's Method)



How it works: Learns word meanings from context (neighboring words).

Example:

If the sentence is "The King wears a crown", Word2Vec learns that:

- ✓ "King" and "Queen" are similar.
- ✓ "King" and "Apple" are not similar.

Problem: Cannot handle **misspelled words** or **new words** that were not in training.



4. GloVe (Stanford's Method)

How it works: Similar to Word2Vec but focuses on word relationships across large texts.



It understands that:

- "Paris" → "France"
- "Berlin" → "Germany"
- "Tokyo" → "Japan"

Problem: Still **static** (one word always has the same meaning).



5. FastText (Facebook's Method)

How it works: Unlike Word2Vec and GloVe, **FastText breaks words into smaller parts** (subwords).

Example:

If a new word "playings" appears, FastText understands it because it has seen:

- "play"
- "playing"
- "ing"

Advantage: Can understand misspelled words and new words.

Problem: Slightly slower than Word2Vec.

6. BERT (Google's Transformer Model - Used in ChatGPT & Google Search)



How it works: Unlike all previous methods, **BERT understands words in context** by reading the entire sentence **in both directions** (before and after a word).

Example:

The word "bank" has two meanings:

"I deposited money in the bank." (Bank = Financial Institution)

"The boat is near the river bank." (Bank = Side of a river)

Old embeddings (Word2Vec, GloVe): Would treat both as the same word.

BERT: Knows they are different based on the sentence.

Best for: Google Search, AI chatbots (ChatGPT), and translation.

Problem: Requires a lot of **computing power**.



8. Deep Learning for Speech Processing

Speech processing means teaching computers to **understand and generate human speech**lt is used in:

- Speech-to-Text (Google Voice Typing)
- Voice Assistants (Siri, Alexa)
- Language Translation (Google Translate)



How Deep Learning Helps in Speech Processing

- 1. Convolutional Neural Networks (CNNs) → Used for Speech Recognition
 - Converts speech into **spectrograms** (like images) and extracts key patterns.
 - Example: Google Voice Assistant uses CNNs to recognize words.
- 2. Recurrent Neural Networks (RNNs) \rightarrow Used for Speech Understanding
 - Reads speech as a sequence of words.
 - **Example:** Google Translate converts text to speech using RNNs.
- 3. Long Short-Term Memory (LSTM) \rightarrow Used for Better Speech Context



- Remembers previous words in a conversation.
- **Example:** Siri remembers your last question for better replies.
- 4. Transformers (BERT, Whisper, Wav2Vec) → Used for Real-Time Speech Translation
 - Processes long conversations faster and more accurately.
 - **Example:** Google Translate voice feature.



9. Explain Photo Encoders and Transfer Learning

A **Photo Encoder** is a type of **neural network** that converts an image into a **compressed numerical representation** (feature vector). This allows the computer to understand and process images efficiently.

How It Works:

- 1. **Input:** A photo is given to the model.
- 2. Feature Extraction: The model detects edges, colors, textures, and objects.
- 3. **Compressed Representation:** The image is converted into a **smaller set of numbers** that represent its features.
- Example:
- In Google Photos, when you search for "dog," it finds dog images because they are encoded into feature vectors.
- In Face Recognition, a person's face is encoded into a unique numerical form for identification.

How Are Photo Encoders Used?

- ✓ Image Search → Google Photos, Pinterest Visual Search
- ✓ Face Recognition → Facebook auto-tagging, Phone Face Unlock
- ✓ **Object Detection** → Self-driving cars identifying pedestrians
- ✓ Medical Imaging → Al detecting diseases from X-rays

What is Transfer Learning?



Transfer Learning is a deep learning technique where a **pre-trained model** is used for a **new task** instead of training a model from scratch.

Why is Transfer Learning Useful?

- Saves time and computing power.
- · Works even with small datasets.
- Learns faster since it already knows general features.

Example:

- A model trained on millions of animal images can quickly learn to classify cats vs. dogs with a small dataset.
- Instead of training a new model, you "transfer" the knowledge of the old model.

How Transfer Learning Works in Images?

- 1. Start with a Pre-Trained Model → Example: ResNet, VGG, MobileNet.
- Remove the Last Layer → The model is already trained on general features like edges and shapes.
- 3. Add New Layers for the New Task → If the original model classified 1,000 objects, modify it to classify only cats and dogs.
- 4. **Fine-Tune the Model** → Train only the last layers while keeping the rest frozen.

Example:

If a model was trained on ImageNet (a large dataset with 1,000 categories), you can
modify it to classify flowers using just a few flower images.

Where Is Transfer Learning Used?

- ✓ **Self-Driving Cars** → Uses models trained on millions of road images.
- ✓ Face Recognition → Uses models pre-trained on celebrity faces to recognize new faces.
- ✓ **Defect Detection in Factories** → Al checks product quality using transfer learning.